

Dossier à rendre

1. TD :	3
2. L'analyse de la compréhension des outils du cloud et de son utilisation	4
3. Expliquer la sécurité votre VM (par le CLOUD)	4
4. Les sources du TP avec des commentaires	4
5. Des procédures d'import des données (voir les fichier sh fournis sur la plateforme et explicité dans les fichiers pdf)	4
6. Vos recommandations par rapport à l'utilisation de ce cloud privé	4

1. TD:

```
root@hadoop-master:~# ls
hadoop-streaming-2.7.2.jar  hbase_create.sh  hbase_odbc_rest.sh  input  mapper.py  purchases2.txt  run-wordcount.sh  setup.sh  start-kafka-zookeeper.sh
happybase.sh               hbase_drop.sh    hdfs              job01.sh  purchases.txt  reducer.py      services_hbase_thrift.sh  start-hadoop.sh  word.txt
```

```
root@hadoop-master:~# ./job01.sh
```

```
Starting namenodes on [hadoop-master]
hadoop-master: Warning: Permanently added 'hadoop-master,172.18.0.2' (ECDSA) to the list of known hosts.
hadoop-master: namenode running as process 273. Stop it first.
hadoop-slave2: Warning: Permanently added 'hadoop-slave2,172.18.0.4' (ECDSA) to the list of known hosts.
hadoop-slave1: Warning: Permanently added 'hadoop-slave1,172.18.0.3' (ECDSA) to the list of known hosts.
hadoop-slave2: datanode running as process 856. Stop it first.
hadoop-slave1: datanode running as process 856. Stop it first.
Starting secondary namenodes [0.0.0.0]
0.0.0.0: secondarynamenode running as process 482. Stop it first.

starting yarn daemons
resourcemanager running as process 666. Stop it first.
hadoop-slave1: Warning: Permanently added 'hadoop-slave1,172.18.0.3' (ECDSA) to the list of known hosts.
hadoop-slave2: Warning: Permanently added 'hadoop-slave2,172.18.0.4' (ECDSA) to the list of known hosts.
hadoop-slave1: nodemanager running as process 970. Stop it first.
hadoop-slave2: nodemanager running as process 970. Stop it first.

hadoop-master: Warning: Permanently added 'hadoop-master,172.18.0.2' (ECDSA) to the list of known hosts.
hadoop-master: zookeeper running as process 15024. Stop it first.
master running as process 15105. Stop it first.
regionserver running as process 15248. Stop it first.
thrifth running as process 12683. Stop it first.
put: 'input/word.txt': File exists
24/05/16 13:57:47 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted outputjob01
```

```
24/05/16 13:57:48 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [mapper.py, reducer.py, /tmp/hadoop-unjar2367100572621533090/] [] /tmp/streamjob2591375481477117521.jar tmpDir=null
24/05/16 13:57:48 INFO client.RMProxy: Connecting to ResourceManager at hadoop-master/172.18.0.2:8032
24/05/16 13:57:49 INFO client.RMProxy: Connecting to ResourceManager at hadoop-master/172.18.0.2:8032
24/05/16 13:57:49 INFO mapred.FileInputFormat: Total input paths to process : 1
24/05/16 13:57:49 INFO mapreduce.JobSubmitter: number of splits:2
24/05/16 13:57:49 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1715847096971_0004
24/05/16 13:57:49 INFO impl.YarnClientImpl: Submitted application application_1715847096971_0004
24/05/16 13:57:49 INFO mapreduce.Job: The url to track the job: http://hadoop-master:8088/proxy/application_1715847096971_0004/
24/05/16 13:57:49 INFO mapreduce.Job: Running job: job_1715847096971_0004
24/05/16 13:57:54 INFO mapreduce.Job: Job job_1715847096971_0004 running in uber mode : false
24/05/16 13:57:54 INFO mapreduce.Job: map 0% reduce 0%
24/05/16 13:57:59 INFO mapreduce.Job: map 100% reduce 0%
24/05/16 13:58:06 INFO mapreduce.Job: map 100% reduce 100%
24/05/16 13:58:06 INFO mapreduce.Job: Job job_1715847096971_0004 completed successfully
24/05/16 13:58:06 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=108
  FILE: Number of bytes written=362015
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=306
  HDFS: Number of bytes written=31
  HDFS: Number of read operations=9
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=4519
  Total time spent by all reduces in occupied slots (ms)=3522
  Total time spent by all map tasks (ms)=4519
  Total time spent by all reduce tasks (ms)=3522
  Total vcore-milliseonds taken by all map tasks=4519
  Total vcore-milliseonds taken by all reduce tasks=3522
  Total megabyte-milliseonds taken by all map tasks=4627456
  Total megabyte-milliseonds taken by all reduce tasks=3606528
Map-Reduce Framework
  Map input records=5
  Map output records=8
  Map output bytes=86
  Map output materialized bytes=114
  Input split bytes=204
  Combine input records=0
  Combine output records=0
  Reduce input groups=3
  Reduce shuffle bytes=114
  Reduce input records=8
  Reduce output records=3
  Spilled Records=16
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=236
  CPU time spent (ms)=1620
  Physical memory (bytes) snapshot=700334000
  Virtual memory (bytes) snapshot=5834657792
  Total committed heap usage (bytes)=521666560
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=192
File Output Format Counters
  Bytes Written=31
24/05/16 13:58:06 INFO streaming.StreamJob: Output directory: outputjob01
```

```
root@hadoop-master:~# hdfs dfs -cat outputjob01/part-00000
Hadoop! 2
Hello   4
Wordcount! 2
```

2. L'analyse de la compréhension des outils du cloud et de son utilisation

Pour l'accès au cloud hidora, j'utilise filezilla pour le transfert de fichier et ensuite grâce à un accès SSH je peux me connecter directement à la machine virtuel.

J'utilise l'endpoint (port 22) pour la connexion.

3. Expliquer la sécurité votre VM (par le CLOUD)

Les différents endpoints avec des ip privés ainsi qu'une identification (login/password) permettent une sécurité du cloud.

4. Les sources du TP avec des commentaires

```
1  #!/usr/bin/env python
2  """mapper.py"""
3  import csv
4  import sys
5
6  # with open('dataaw_fro03.csv', newline='') as csvfile:
7  csv_reader = csv.reader(sys.stdin)
8
9  next(csv_reader, None)
10
11
12  for row in csv_reader:
13      nomcli, prenomcli, cpcli, villecli, datcde, timbrecli, qte = (
14          row[2],
15          row[3],
16          row[4],
17          row[5],
18          row[7],
19          row[9],
20          row[15],
21      )
22      cpcli = int(cpcli) / 1000
23      if cpcli in [53, 61, 28] and '2008' < datcde < '2016':
24          client = nomcli + ' ' + prenomcli
25          dep = cpcli
26
27          print(str(timbrecli) + ';' + str(client) + ';' + str(villecli) + ';' + str(dep) + ';' + str(qte))
```

On est sur un mapper assez simple utilisant la librairie csv pour lire le fichier csv ouvert par hadoop.

Ensuite on identifie les différentes colonnes qui nous intéresse.

On filtre les données qui ne nous interesse pas et on les envoie pour le traitement par le reducer.

```

1 > import ...
5
6 results = {}
7 resultsdep = {}
8
9 for line in sys.stdin:
10     line = line.strip()
11     timbrecli, client, villecli, dep, qte = line.split(";")
12
13     try:
14         qte = int(qte)
15     except ValueError:
16         continue
17
18     # Créer une clé unique pour chaque (client, ville, département)
19     cle = (client, villecli, dep)
20
21     # Mettre à jour le dictionnaire en incrémentant le compteur pour cette clé
22     if cle in results:
23         results[cle] += qte
24     else:
25         results[cle] = qte
26
27     # Mettre à jour le dictionnaire des résultats par département
28     if dep in resultsdep:
29         resultsdep[dep] += qte
30     else:
31         resultsdep[dep] = qte
32
33 # Création d'un DataFrame à partir de results
34 results_df = pd.DataFrame(list(results.items()), columns=['Client_Ville_Dep', 'Total_Commandes'])
35
36 # Export des résultats vers un fichier Excel
37 results_excel_file = '/datavolume1/resultats.xlsx'
38 results_df.to_excel(results_excel_file, index=False)
39
40 # Créer un DataFrame à partir du dictionnaire des résultats par département
41 df = pd.DataFrame(list(resultsdep.items()), columns=['Département', 'Total'])
42
43 # Créer une nouvelle figure
44 plt.figure()
45
46 # Créer le graphe pie
47 plt.pie(df['Total'], labels=df['Département'], autopct='%1.1f%%', startangle=140)
48 plt.axis('equal') # Assurer que le diagramme est circulaire
49 plt.title("Répartition des commandes par département")
50
51 # Enregistrer le graphe au format PDF
52 output_pdf_file = '/datavolume1/resultat.pdf' # Chemin où le fichier PDF sera enregistré
53 with PdfPages(output_pdf_file) as pdf:
54     pdf.savefig() # Sauvegarder le graphe dans le fichier PDF
55

```

Le reducer traite les différentes données en deux objets, transforme ensuite en dataframe.

Le premier permet d'afficher dans un excel les clients et la quantités de commandes qu'ils ont effectués.

Le deuxième les commandes par départements affiché dans une pie faites avec matplotlib.

5. Des procédures d'import des données (voir les fichier sh fournis sur la plateforme et explicité dans les fichiers pdf)

Pour pouvoir ensuite lancer mon mapper et reducer, j'adapte le fichier job01.sh.

```
1 ▶ cp /usr/local/hadoop/share/hadoop/tools/lib/hadoop-streaming-2.7.2.jar .
2   ./start-hadoop.sh
3   start-hbase.sh
4   hbase-daemon.sh start thrift
5   hdfs dfs -mkdir -p input
6   hdfs dfs -put dataw_fro83.csv input
7   hdfs dfs -rm -r outputjob82
8   hadoop jar hadoop-streaming-2.7.2.jar -file mapper1.py -mapper "python3 mapper1.py" -file reducer1.py -reducer "python3 reducer1.py" -input input/dataw_fro83.csv -output outputjob82
9
```

Ensuite grâce à Filezilla, j'envoie tous les documents sur le cloud, je m'y connecte en SSH.

Je copie tout sur le docker, et j'exécute le job01.sh.

```
root@hadoop-master:~# ./job01.sh
```

```
Starting namenodes on [hadoop-master]
hadoop-master: Warning: Permanently added 'hadoop-master,172.18.0.2' (ECDSA) to the list of known hosts.
hadoop-master: namenode running as process 273. Stop it first.
hadoop-slave2: Warning: Permanently added 'hadoop-slave2,172.18.0.4' (ECDSA) to the list of known hosts.
hadoop-slave1: Warning: Permanently added 'hadoop-slave1,172.18.0.3' (ECDSA) to the list of known hosts.
hadoop-slave2: datanode running as process 856. Stop it first.
hadoop-slave1: datanode running as process 856. Stop it first.
Starting secondary namenodes [0.0.0.0]
0.0.0.0: secondarynamenode running as process 482. Stop it first.

starting yarn daemons
resourcemanager running as process 666. Stop it first.
hadoop-slave2: Warning: Permanently added 'hadoop-slave2,172.18.0.4' (ECDSA) to the list of known hosts.
hadoop-slave1: Warning: Permanently added 'hadoop-slave1,172.18.0.3' (ECDSA) to the list of known hosts.
hadoop-slave2: nodemanager running as process 970. Stop it first.
hadoop-slave1: nodemanager running as process 970. Stop it first.

hadoop-master: Warning: Permanently added 'hadoop-master,172.18.0.2' (ECDSA) to the list of known hosts.
hadoop-master: zookeeper running as process 15024. Stop it first.
master running as process 15105. Stop it first.
: regionserver running as process 15248. Stop it first.
thrift running as process 12683. Stop it first.
put: 'input/dataaw_fro03.csv': File exists
24/05/20 14:04:07 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted outputjob02
24/05/20 14:04:08 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [mapper1.py, reducer1.py, /tmp/hadoop-unjar8437968220122313420/] [] /tmp/streamjob5349489168750919770.jar tmpDir=null
24/05/20 14:04:09 INFO client.RMProxy: Connecting to ResourceManager at hadoop-master/172.18.0.2:8032
24/05/20 14:04:09 INFO client.RMProxy: Connecting to ResourceManager at hadoop-master/172.18.0.2:8032
24/05/20 14:04:09 INFO mapred.FileInputFormat: Total input paths to process : 1
24/05/20 14:04:09 INFO mapreduce.JobSubmitter: number of splits:2
24/05/20 14:04:09 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1715847096971_0019
24/05/20 14:04:09 INFO impl.YarnClientImpl: Submitted application application_1715847096971_0019
24/05/20 14:04:09 INFO mapreduce.Job: The url to track the job: http://hadoop-master:8088/proxy/application_1715847096971_0019/
24/05/20 14:04:09 INFO mapreduce.Job: Running job: job_1715847096971_0019
24/05/20 14:04:15 INFO mapreduce.Job: Job job_1715847096971_0019 running in uber mode : false
24/05/20 14:04:15 INFO mapreduce.Job: map 0% reduce 0%
24/05/20 14:04:19 INFO mapreduce.Job: map 100% reduce 0%
24/05/20 14:04:25 INFO mapreduce.Job: map 100% reduce 100%
24/05/20 14:04:25 INFO mapreduce.Job: Job job_1715847096971_0019 completed successfully
24/05/20 14:04:25 INFO mapreduce.Job: Counters: 50
```

```
24/05/20 14:04:25 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=67712
    FILE: Number of bytes written=497274
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=26536758
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=9
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Killed map tasks=1
    Launched map tasks=2
    Launched reduce tasks=1
    Data-local map tasks=2
    Total time spent by all maps in occupied slots (ms)=5377
    Total time spent by all reduces in occupied slots (ms)=2875
    Total time spent by all map tasks (ms)=5377
    Total time spent by all reduce tasks (ms)=2875
    Total vcore-milliseconds taken by all map tasks=5377
    Total vcore-milliseconds taken by all reduce tasks=2875
    Total megabyte-milliseconds taken by all map tasks=5506048
    Total megabyte-milliseconds taken by all reduce tasks=2944000
  Map-Reduce Framework
    Map input records=135278
    Map output records=1657
    Map output bytes=64392
    Map output materialized bytes=67718
    Input split bytes=218
    Combine input records=0
    Combine output records=0
    Reduce input groups=949
    Reduce shuffle bytes=67718
    Reduce input records=1657
    Reduce output records=0
    Spilled Records=3314
    Shuffled Maps =2
    Failed Shuffles=0
    Merged Map outputs=2
    GC time elapsed (ms)=228
    CPU time spent (ms)=2500
    Physical memory (bytes) snapshot=701792256
    Virtual memory (bytes) snapshot=5832609792
    Total committed heap usage (bytes)=526385152
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=26536540
  File Output Format Counters
    Bytes Written=0
24/05/20 14:04:25 INFO streaming.StreamJob: Output directory: outputjob02
```


Tout est alors envoyé dans le datavolume1 du docker.

Je les récupères grâce à Filezilla et les mets sur ma machine locale.

6. Vos recommandations par rapport à l'utilisation de ce cloud privé

Je n'avais encore jamais utilisé de cloud pour le travail de cette manière et j'ai trouvé cela très pratique. Mon mac (avec une puce M1) ne me permettant pas de pouvoir installé Hadoop sur ma machine, cela permet donc de mettre en place un espace de travail totalement géré et indépendant de la machine de chaque utilisateur.